

Class 2: Moving averages and ARMA

Andrew Parnell
andrew.parnell@ucd.ie



Learning outcomes

- ▶ Recognise and understand the basic theory behind MA(1) and MA(q) models
- ▶ Understand the basic ARMA(p,q) formulation
- ▶ Know the basics of using the forecast package
- ▶ Understand the limitations of ARMA forecasting

1 / 27

2 / 27

Reminder: The most important slide in the course

Almost all of time series is based on two ideas:

1. Base your future predictions on previous values of the data
2. **Base your future predictions on how wrong you were in your past predictions**

3 / 27

Reminder: AR models

- ▶ An Autoregressive (AR) model works by making the current data point dependent on the previous value, dampened by a parameter
- ▶ The usual likelihood used is:

$$y_t \sim N(\alpha + \beta y_{t-1}, \sigma^2)$$

- ▶ β is usually constrained (naturally via the data) to lie between -1 and 1. Outside that range the process blows up
- ▶ The sample PACF is often a good way of diagnosing if an AR model might be appropriate

4 / 27

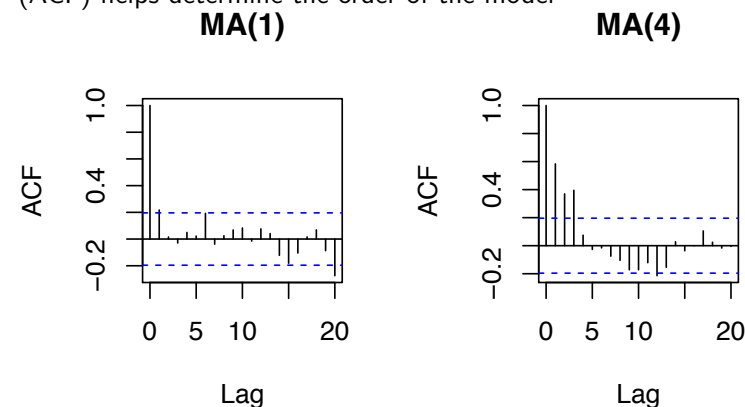
Intro to Moving Average Models

- ▶ Moving Average (MA) models are similar to AR models but they depend on the previous residual of the series rather than the value itself
- ▶ The previous residual is made up of how well we forecasted the last value of the series
- ▶ If the previous residual was large (i.e. our forecast was bad) then we want to make a big change to the next prediction
- ▶ If the previous residual was small (i.e. our forecast was good) then we might not want to make much of a change

5 / 27

Moving average models and the ACF/PACF

- ▶ Recall that the sample partial autocorrelation function (PACF) can be used to diagnose whether an AR model is appropriate (and also suggest the order p)
- ▶ For the MA model, it is the sample autocorrelation function (ACF) helps determine the order of the model



6 / 27

Example 1: MA(1)

- ▶ The MA(1) model is defined as:

$$y_t = \alpha + \theta \epsilon_{t-1} + \epsilon_t$$

where $\epsilon_t \sim N(0, \sigma^2)$ as usual

- ▶ Parameter α represents the overall mean, whilst θ controls the amount of weight placed on previous residuals
- ▶ Like the AR model the values of θ are not expected to be outside $(-1, 1)$, and negative values can sometimes be physically unrealistic
- ▶ The likelihood version of the model is:

$$y_t \sim N(\alpha + \theta \epsilon_{t-1}, \sigma^2)$$

7 / 27

Simulating from the MA(1) process


Below is some simple code to simulate from an MA(1) process. Note that the first values of y and ϵ need to be initialised

```
T = 100 # Number of observations
sigma = 1 # Residual sd
alpha = 0 # Mean
theta = runif(1) # Choose a positive value
y = eps = rep(NA, T)
y[1] = alpha
eps[1] = 0
for(t in 2:T) {
  y[t] = rnorm(1, mean = alpha + theta * eps[t-1],
              sd = sigma)
  eps[t] = y[t] - alpha - theta * eps[t-1]
}
```

8 / 27

Time series plot

```
plot(1:T,y,type='l')
```



9 / 27

Fitting MA(1) models

- We can fit an MA(1) model with the forecast package like before

```
Arima(y, order = c(0, 0, 1))
```

```
## Series: y
## ARIMA(0,0,1) with non-zero mean
##
## Coefficients:
##          ma1      mean
##          0.2118 -0.0325
## s.e.    0.0924   0.1120
##
## sigma^2 estimated as 0.8747:  log likelihood=-134.21
## AIC=274.43   AICc=274.68   BIC=282.24
```

10 / 27

Extending to MA(q)

- As with the AR(p) process we can extend this model to have the current value of y depending on more than one previous residual
- The model becomes an MA(q) model with:

$$y_t \sim N(\alpha + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \dots + \theta_q \epsilon_{t-q}, \sigma^2)$$

- The parameters are as before, except there are now q values of θ .
- Usually when estimated they will decrease with q ; the older residuals matter less

11 / 27

Fitting an MA(q) model

```
Arima(y, order = c(0, 0, 3))
```

```
## Series: y
## ARIMA(0,0,3) with non-zero mean
##
## Coefficients:
##          ma1      ma2      ma3      mean
##          0.2269  0.0170 -0.0413 -0.0338
## s.e.    0.1005  0.0981  0.0876  0.1111
##
## sigma^2 estimated as 0.8897:  log likelihood=-134.04
## AIC=278.07   AICc=278.71   BIC=291.1
```

- Compare the AIC of this model with the previous MA(1) version

12 / 27

Forecasting an MA value

- ▶ You can create a one step ahead forecast for an MA(1) model by:

$$\hat{y}_{t+1} = \alpha + \theta\epsilon_t$$

- ▶ Forecasts of more than one step ahead will be pretty boring, as every future prediction of $\hat{\epsilon}_t$ will be 0
- ▶ Thus MA(q) models are only really informative if you are forecasting $q - 1$ steps ahead

13 / 27

Combining AR and MA into ARMA

- ▶ There is no reason why we have to use just AR or MA on their own
- ▶ It's possible to combine them together, for example:

$$y_t = \alpha + \beta y_{t-1} + \theta\epsilon_{t-1} + \epsilon_t$$

This is an *Autoregressive Moving Average* (ARMA) model

- ▶ It's often written as ARMA(p,q) where p is the number of AR terms (here 1) and q the number of MA terms (here also 1)
- ▶ ARMA models can deal with a very wide variety of flexible time series behaviour, though they remain stationary
- ▶ The likelihood format is:

$$y_t \sim N(\alpha + \beta y_{t-1} + \theta\epsilon_{t-1}, \sigma^2)$$

14 / 27

Fitting an ARMA(1, 1) model

```
Arima(y, order = c(1, 0, 1))
```

```
## Series: y
## ARIMA(1,0,1) with non-zero mean
##
## Coefficients:
##          ar1      ma1      mean
##      0.1119  0.1085  -0.0323
## s.e.  0.3743  0.3703  0.1153
##
## sigma^2 estimated as 0.8829:  log likelihood=-134.17
## AIC=276.33   AICc=276.76   BIC=286.76
```

- ▶ Compare again with previous models

15 / 27

The general ARMA(p, q) framework

- ▶ The general equation for an ARMA(p, q) model is:

$$y_t = \alpha + \sum_{i=1}^p \beta_i y_{t-i} + \sum_{j=1}^q \theta_j \epsilon_{t-j} + \epsilon_t$$

- ▶ The values of β have to be tightly controlled to get a series that is stationary, though this is only really a problem if we want to simulate the time series
- ▶ Occasionally you will run into problems with Arima because it doesn't use maximum likelihood (by default) to fit the models. It uses something faster and more approximate instead

16 / 27


Predicting the future with ARMA

- ▶ The forecast package contains methods to predict into the future
- ▶ First create a model (here ARMA(2, 1))

```
my_model = Arima(y, order = c(2, 0, 1))
```

- ▶ ...then forecast...

```
plot(forecast(my_model, h = 10))
```



17 / 27

A real-world example

Steps in a time series analysis


1. Plot the data and the ACF/PACF
2. Decide if the data look stationary or not. If not, perform a suitable transformation and return to 1
3. Guess at a suitable p and q for an ARMA(p , q) model
4. Fit the model
5. Try a few models around it by increasing/decreasing p and q and checking the AIC (or others)
6. Check the residuals
7. Forecast into the future

19 / 27

A real example: wheat data

- ▶ Let's follow the steps for the wheat data:

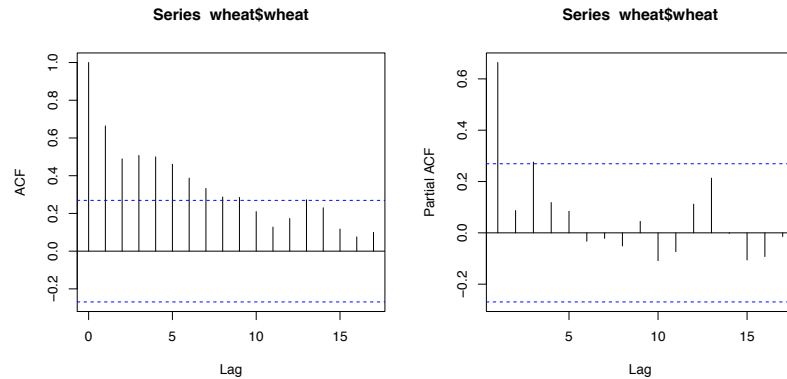
```
wheat = read.csv('../data/wheat.csv')  
plot(wheat$year, wheat$wheat, type = 'l')
```



20 / 27

ACF and PACF

```
par(mfrow = c(1, 2))
acf(wheat$wheat)
pacf(wheat$wheat)
```



- Suggest starting with AR(1) or AR(3)?

21 / 27

First model

```
Arima(wheat$wheat, order = c(1, 0, 0))
```

```
## Series: wheat$wheat
## ARIMA(1,0,0) with non-zero mean
##
## Coefficients:
##          ar1          mean
##          0.8972 20849.522
## s.e.    0.0826  3615.699
##
## sigma^2 estimated as 10079564: log likelihood=-502.34
## AIC=1010.68   AICc=1011.17   BIC=1016.59
```

22 / 27

Next models

- Try AR(2), ARMA(1, 1), and ARMA(2, 1)

```
Arima(wheat$wheat, order = c(2, 0, 0))$aic
```

```
## [1] 1012.683
```

```
Arima(wheat$wheat, order = c(1, 0, 1))$aic
```

```
## [1] 1011.36
```

```
Arima(wheat$wheat, order = c(2, 0, 1))$aic
```

```
## [1] 1004.125
```


- Best one seems to be ARMA(2, 1). (could also try others)

23 / 27

Check residuals

- Check the residuals of this model


```
my_model_ARMA21 = Arima(wheat$wheat, order = c(2, 0, 1))
qqnorm(my_model_ARMA21$residuals)
qqline(my_model_ARMA21$residuals)
```



24 / 27

Forecast into the future

```
plot(forecast(my_model_ARMA21,h=20))
```



What happened to the forecasts here?

- ▶ Why did the series diverge rapidly away from what you might have expected?
- ▶ The answer is that we have fitted a *stationary model*, i.e. one with constant mean and variance
- ▶ The model will just slowly revert back to that mean over time. The speed at which it reverts will depend on the amount of autocorrelation in the series
- ▶ The solution to this lies in better identification of the trend. See the next lecture!

26 / 27

Summary

- ▶ MA(q) models are used to create future forecasts based on the error in the previous forecasts
- ▶ ARMA models combine AR and MA ideas together
- ▶ The forecast package allows us to fit all of these models
- ▶ We need to be a bit careful with forecasts that assume stationarity - they will mean-revert