Class 3: Integrated models and ARIMA

Andrew Parnell
andrew.parnell@ucd.ie

Learning outcomes

- ► Understand how differencing works to help make data stationary
- ► Know the basics of the ARIMA(p, d, q) framework
- ► Understand how to fit an ARIMA(p, d, q) model in a realistic setting

Reminder: stationarity

- A time series is said to be weakly stationary if:
    - The mean is stable
    - The variance is stable
    - The autocorrelation doesn't depend on where you are in the series

Reminder: ARMA models

- Combine the autoregressive and the moving average framework into one

- The general equation for an ARMA(p, q) model is:

$$y_t = \alpha + \sum_{i=1}^{p} \beta_i y_{t-i} + \sum_{j=1}^{q} \theta_j \epsilon_{t-q} + \epsilon_t$$

Combining ARMA with the random walk to produce ARIMA

▶ There is one other time series model we have already met, that of the random walk:

$$y_t = y_{t-1} + \epsilon_t$$

where $\epsilon_t \sim N(0, \sigma^2)$

▶ We could re-write this as:

$$y_t - y_{t-1} = \epsilon_t$$

i.e. the *differences* are random normally-distributed noise

Differencing

- Differencing is a great way of getting rid of a linear trend
- If $y_t \approx y_{t-1} + b$ then there will be an increasing linear slope in the time series.
- Creating $y_t - y_{t-1}$ will remove it and all values will hover around the value b
- Differencing twice will remove a quadratic trend for the same reasons
- You can do even higher levels of differencing but this starts to cause problems
- The twice differenced series is:

$$(y_t - y_{t-1}) - (y_{t-1} - y_{t-2}) = y_t - 2y_{t-1} + y_{t-2}$$

Idea: combine differencing into the ARMA framework

- We can combine these ideas into the ARMA framework to produce an ARIMA model (the I stands for integrated, i.e. differenced)
- An ARIMA model isn't really stationary as the differences are actually removing part of the trend
- The ARIMA model is written as ARIMA(p,d,q) where p and q are as before and d is the number of differences.

Example: the ARIMA(1,1,1) model

- If we want to fit an ARIMA(1,1,1) model we first let $z_t = y_t - y_{t-1}$ then fit the model:

$$z_t \sim N(\alpha + \beta z_{t-1} + \theta \epsilon_{t-1}, \sigma^2)$$

- This is equivalent to an ARMA model on the first differences

Fitting an ARIMA(1, 1, 1) model to the wheat data

- Recall that the ARMA(2,1) fit wasn't very good to the wheat data
- Instead try an ARIMA(1, 1, 0) model (i.e. AR(1) on the first differences)

```
wheat = read.csv('../../data/wheat.csv')
Arima(wheat$wheat, order = c(1, 1, 0))
```

```
## Series: wheat$wheat
## ARIMA(1,1,0)
##
## Coefficients:
##          ar1
##       -0.0529
## s.e.   0.1520
##
## sigma^2 estimated as 10076177:  log likelihood=-492.55
## AIC=989.1   AICc=989.34   BIC=993
```

General format: the ARIMA(p,d,q) model

- First take the dth difference of the series y_t, and call this z_t.
- If you want to do this by hand in R you can use the `diff` function, e.g. `diff(y, differences = 2)`
- Then fit the model:

$$z_t \sim N\left(\alpha + \sum_{i=1}^{p} \beta_i z_{t-i} + \sum_{j=1}^{q} \theta_j \epsilon_{t-j}, \sigma^2\right)$$

Choosing p, d and q

- It's much harder to have an initial guess at all of p, d and q in one go.
- We can usually guess at the number of differences d from the time series and ACF plots. If there is a very high degree of autocorrelation it's usually a good idea to try a model with $d=1$ or 2
- I've never met a model where you needed to difference more than twice. Beware of over-differencing

Revisiting the real-world example

Steps in an ARIMA time series analysis

1. Plot the data and the ACF/PACF
2. Decide if the data look stationary or not. If not, perform a suitable transformation and return to 1. **If the data has a strong trend or there is a high degree of autocorrelation try 1 or 2 differences**
3. Guess at values of p, d, and q for an ARIMA(p, d, q) model
4. Fit the model
5. Try a few models around it by increasing/decreasing p, d and q and checking the AIC (or others)
6. Check the residuals
7. Forecast into the future

A real example: wheat data

- Plot reminder

```
wheat = read.csv('../../data/wheat.csv')
plot(wheat$year, wheat$wheat, type = 'l')
```

ACF and PACF

```r
par(mfrow = c(1, 2))
acf(wheat$wheat)
pacf(wheat$wheat)
```


- Suggest looking at first differences

Plot of first differences

```
plot(diff(wheat$wheat), type = 'l')
```

ACF/PACF of first differences

```
par(mfrow = c(1, 2))
acf(diff(wheat$wheat))
pacf(diff(wheat$wheat))
```


Series diff(wheat$wheat)   **Series diff(wheat$wheat)**

▶ Interesting peaks in ACF at lag 2, and PACF at lag 2.

First model

```
Arima(wheat$wheat, order = c(0, 1, 0))
```

```
## Series: wheat$wheat
## ARIMA(0,1,0)
##
## sigma^2 estimated as 9905911:  log likelihood=-492.61
## AIC=987.22   AICc=987.3   BIC=989.17
```

Next models

- ▶ Try ARIMA(1, 1, 1), ARIMA(1, 1, 0), ARIMA(0, 1, 1)

```
Arima(wheat$wheat, order = c(1, 1, 1))$aic
```

[1] 985.4555

```
Arima(wheat$wheat, order = c(1, 1, 0))$aic
```

[1] 989.0983

```
Arima(wheat$wheat, order = c(0, 1, 1))$aic
```

[1] 986.0402

- ▶ Best one seems to be ARIMA(1, 1, 1). (though BIC suggests others)

Check residuals

- ▶ Check the residuals of this model

```
my_model_ARIMA111 = Arima(wheat$wheat, order = c(1, 1, 1))
qqnorm(my_model_ARIMA111$residuals)
qqline(my_model_ARIMA111$residuals)
```

Normal Q–Q Plot

Forecast into the future

```r
plot(forecast(my_model_ARIMA111,h=20))
```

Forecasts from ARIMA(1,1,1)

Why does the difference not continus into the future?

- You might have expected the forecasts to continue rising into the future due to the difference
- The MA part of the model is obviously flat as previously discussed as there are no further errors to correct
- The AR part of the model reverts back to the estimated mean of the last data point because the β parameter is less than 1 - it dampens out the future predictions and stops them from going crazy

Summary

- ARIMA models extend the ARMA framework to further add in differencing
- A single difference will remove a linear trend, a second difference a quadratic trend
- Can spot the need for differencing from the time series plot and the ACF
- Do not over-difference your data!